

Guidance to run PyTorch BERT-Large PreTraining on NVIDIA H100 GPUs

Login to ACES cluster and run the commands below.

```
$cd $SCRATCH
```

```
$mkdir h100-benchmarks
```

```
$cd h100-benchmarks
```

```
$git clone -b core_r0.4.0 https://github.com/NVIDIA/Megatron-LM.git
```

```
# Below change is required to print average throughput
```

```
# Update $SCRATCH/Megatron-LM/megatron/training.py with (utils/training.py)
```

```
# create a slurm job file test_pytorch_bert_large.slurm and copy and paste the content below to it.
```

```
$vim test_pytorch_bert_large.slurm
```

```
#!/bin/bash
```

```
##ESSARY JOB SPECIFICATIONS
```

```
#SBATCH --job-name=<your_job_name>
```

```
#SBATCH --time=1:00:00          #Set the wall clock limit to 5hr
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks=1
```

```
#SBATCH --cpus-per-task=8
```

```
#SBATCH --mem=80GB
```

```
#SBATCH --output=<your_job>_run.%j
```

```
#SBATCH --partition=gpu          #Request 1 GPU per node can be 1 or
```

```
#SBATCH --gres=gpu:h100:1       #Request 1 GPU per node can be 1 or 2
```

```
#This command is used to get stats of H100 GPU utilization
```

```
nvidia-smi
```

```
--query-gpu=timestamp,name,pci.bus_id,driver_version,pstate,pcie.link.gen.max,pcie.link.gen.current,temperature.gpu,utilization.gpu,utilization.memory,memory.total,memory.free,memory.used --format=csv -l 1 > <your_job>_GPU_stats.log &
```

```
#This command is used to get stats of CPU cores utilization  
watch -n 5 ps -u $USER > <your_job>_CPU_stats.log &
```

```
export  
SINGULARITY_BINDPATH="$SCRATCH/h100-benchmarks/Megatron-LM:/workspace,  
/scratch/data/pytorch-language-modelling-datasets:/shared_space_datasets"
```

```
export CUDA_DEVICE_MAX_CONNECTIONS=1
```

```
GPUS_PER_NODE=1  
NNODES=1  
WORLD_SIZE=$((($GPUS_PER_NODE*$NNODES))  
NPROCS_PER_NODE=1  
NPROCS=1
```

```
MASTER=`/bin/hostname -s`  
SLAVES=`scontrol show hostnames $SLURM_JOB_NODELIST | grep -v $MASTER`  
HOSTLIST="$MASTER $SLAVES"  
echo $HOSTLIST
```

```
echo head node: $MASTER
```

```
VOCAB_FILE=/shared_space_datasets/vocab.txt  
DATA_PATH=/shared_space_datasets/intel-bert_text_sentence
```

```
micro_batch_size=64  
global_batch_size=64  
train_iters=5000  
precision=bf16
```

```
BERT_ARGS="  
  --num-layers 24 \  
  --hidden-size 1024 \  
  --num-attention-heads 16 \  
  --seq-length 512 \  
  --max-position-embeddings 512 \  
"
```

```
--micro-batch-size ${micro_batch_size} \  
--global-batch-size ${global_batch_size} \  
--lr 0.0001 \  
--train-iters ${train_iters} \  
--lr-decay-iters 990000 \  
--lr-decay-style linear \  
--min-lr 0.00001 \  
--weight-decay 1e-2 \  
--lr-warmup-fraction .01 \  
--clip-grad 1.0 \  
--${precision}  
"
```

```
DATA_ARGS="  
  --data-path $DATA_PATH \  
  --vocab-file $VOCAB_FILE \  
  --split 450,32,20  
"
```

```
OUTPUT_ARGS="  
  --log-interval 100 \  
  --save-interval 10000 \  
  --eval-interval 1000 \  
  --eval-iters 10  
"
```

```
    singularity exec --nv  
/scratch/data/containers/nvidia-containers/pytorch-nemo-23-06.sif torchrun --nnodes  
$NNODES \  
  --nproc_per_node $NPROCS_PER_NODE \  
  --rdzv_id $RANDOM \  
  --rdzv_backend c10d \  
  --rdzv_endpoint $MASTER:${RANDOM} \  
/workspace/pretrain_bert.py \  
$BERT_ARGS \  
$DATA_ARGS \  
$OUTPUT_ARGS
```

```
$sbatch test_pytorch_bert_large.slurm
```

